

Car Price Prediction using K-Nearest Neighbors (KNN)

AIML ASSIGNMENT REPORT

1CS101CC22-Introduction to AI-ML

Namdar Mohammad.Mehdi
24BEI041
Lakshya Jain
24BEI043

Car Price Prediction using K-Nearest Neighbors (KNN)

1. Introduction

Car price prediction is a significant application in the field of machine learning and data science. The ability to accurately predict the price of a used car based on its features can help buyers and sellers make informed decisions. In this project, we have implemented the K-Nearest Neighbors (KNN) algorithm to predict the selling price of used cars using real-world data.

2. Dataset Description

We have used the CarDekho dataset, which contains various attributes of used cars, including:

- Year – Manufacturing year of the car
- Mileage – Distance the car can travel per liter of fuel
- Engine – Engine capacity in CC
- Power – Horsepower of the car
- Selling Price – Price at which the car is sold (Target Variable)

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Load the dataset
file_path = "cardekho_dataset.csv" # Change if needed
df = pd.read_csv(file_path)

# Display first few rows
print(df.head())
```

Unnamed: 0	car_name	brand	model	vehicle_age	km_driven	\
0	0	Maruti Alto	Maruti Alto	9	120000	
1	1	Hyundai Grand	Hyundai Grand	5	20000	
2	2	Hyundai i20	Hyundai i20	11	60000	
3	3	Maruti Alto	Maruti Alto	9	37000	
4	4	Ford Ecosport	Ford Ecosport	6	30000	

	seller_type	fuel_type	transmission_type	mileage	engine	max_power	seats	\
0	Individual	Petrol	Manual	19.70	796	46.30	5	
1	Individual	Petrol	Manual	18.90	1197	82.00	5	
2	Individual	Petrol	Manual	17.00	1197	80.00	5	
3	Individual	Petrol	Manual	20.92	998	67.10	5	
4	Dealer	Diesel	Manual	22.77	1498	98.59	5	

	selling_price
0	120000
1	550000
2	215000
3	226000
4	570000

3. Data Preprocessing

Before applying our model, we performed the following preprocessing steps:

- Handling missing values
- Removing unnecessary columns (e.g., 'Seats')
- Feature scaling and normalization for better accuracy
- Splitting data into training (80%) and testing (20%)

```
print("Basic Statistical Analysis:\n",  
df.describe())
```

Basic Statistical Analysis:

	Unnamed: 0	vehicle_age	km_driven	mileage	engine
count	15411.000000	15411.000000	1.541100e+04	15411.000000	15411.000000
mean	9811.857699	6.036338	5.561648e+04	19.701151	1486.057751
std	5643.418542	3.013291	5.161855e+04	4.171265	521.106696
min	0.000000	0.000000	1.000000e+02	4.000000	793.000000
25%	4906.500000	4.000000	3.000000e+04	17.000000	1197.000000
50%	9872.000000	6.000000	5.000000e+04	19.670000	1248.000000
75%	14668.500000	8.000000	7.000000e+04	22.700000	1582.000000
max	19543.000000	29.000000	3.800000e+06	33.540000	6592.000000

	max_power	seats	selling_price
count	15411.000000	15411.000000	1.541100e+04
mean	100.588254	5.325482	7.749711e+05
std	42.972979	0.807628	8.941284e+05
min	38.400000	0.000000	4.000000e+04
25%	74.000000	5.000000	3.850000e+05
50%	88.500000	5.000000	5.560000e+05
75%	117.300000	5.000000	8.250000e+05
max	626.000000	9.000000	3.950000e+07

```
# Q-2: Drop Null Records & Encode Categorical Data  
# Drop unnecessary columns  
drop_cols = ['mileage', 'engine', 'max_power', 'seats', 'seller_type', 'fuel_type', 'car_name', 'transmission_type']  
df = df.drop(columns=drop_cols)  
  
# Label Encoding for Categorical Data  
label_encoders = {}  
for col in ['brand', 'model']:  
    le = LabelEncoder()  
    df[col] = le.fit_transform(df[col])  
    label_encoders[col] = le
```

4. Methodology

4.1 Algorithm: K-Nearest Neighbors (KNN)

KNN is a simple, yet powerful, algorithm that predicts values based on the similarity of data points. It works as follows:

1. Select the number of neighbors (k).
2. Compute the distance between the new data point and all other points in the dataset.
3. Select k nearest points.
4. Compute the average price of the selected points (for regression problems).
5. Assign the predicted price.

5. Code Implementation

Training and Testing:

```
# Predict Log prices
y_pred_log = knn.predict(X_test)

# Convert Log predictions back to normal price
y_pred = np.expml(y_pred_log)

print("Predictions made successfully!")

# Calculate Error Metrics
mae = mean_absolute_error(np.expml(y_test), y_pred) #
mse = mean_squared_error(np.expml(y_test), y_pred)
r2 = r2_score(np.expml(y_test), y_pred)
rmse = np.sqrt(mse)
print(f"Mean Absolute Error (MAE): ₹{mae:,.2f}")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"R2 Score: {r2:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")

# Q-5: Train KNN Model
# Train KNN model
knn = KNeighborsRegressor(n_neighbors=2, weights='distance') # Weighted KNN
knn.fit(X_train, y_train)

print("KNN Model Trained Successfully!")
```

Considering Depreciation for a realistic outcome

```
def apply_realistic_depreciation(predicted_price, vehicle_age):
    if vehicle_age == 0:
        return predicted_price # New car, no depreciation
    elif vehicle_age <= 3:
        return predicted_price * (0.75) # 25% depreciation in 3 years
    elif vehicle_age <= 5:
        return predicted_price * (0.60) # 40% depreciation in 5 years
    elif vehicle_age <= 10:
        return predicted_price * (0.40) # 60% depreciation in 10 years
    else:
        return predicted_price * (0.30) # Older cars retain ~30% value

# Modify prediction function to include depreciation
def predict_car_price(brand_name, model_name, vehicle_age, km_driven):
    # Encode brand & model
    if brand_name in label_encoders['brand'].classes_ and model_name in label_encoders['model'].classes_:
        brand_encoded = label_encoders['brand'].transform([brand_name])[0]
        model_encoded = label_encoders['model'].transform([model_name])[0]
    else:
        return "Error: Invalid brand or model name."


    # Prepare input & scale it
    input_data = np.array([[brand_encoded, model_encoded, vehicle_age, km_driven]])
    input_data = scaler.transform(input_data)

    # Predict log price & convert back
    predicted_log_price = knn.predict(input_data)[0]
    predicted_price = np.expml(predicted_log_price)

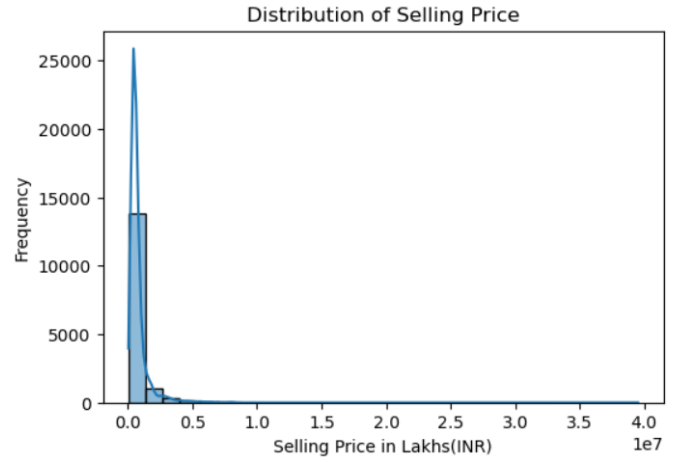
    # Apply improved depreciation formula
    final_price = apply_realistic_depreciation(predicted_price, vehicle_age)
```

6. Statistical Analysis & Visualizations


6.1 VEHICLE AGE VS SELLING PRICE



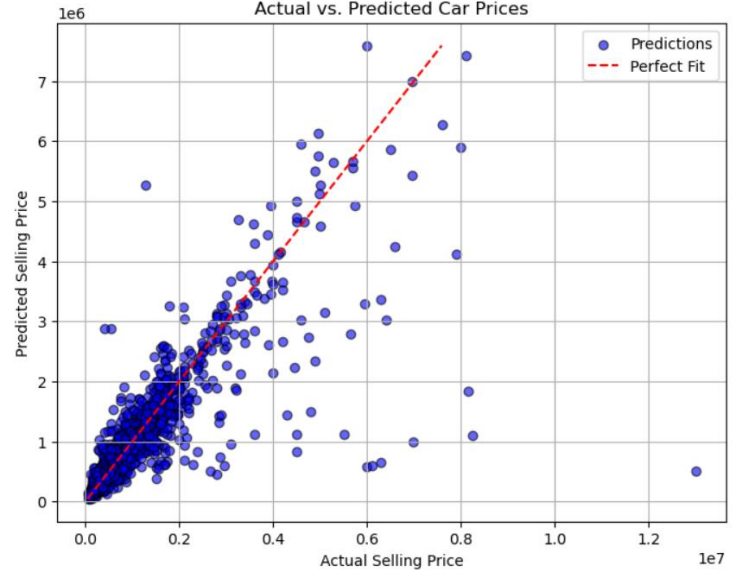
6.2 DISTRIBUTION OF SELLING PRICE



6.3 AVERAGE SELLING PRICE BY AGE



6.4 EXPECTED PRICE VS REAL PRICE



7. Model Evaluation

<i>Metric</i>	<i>Value</i>
MAE	₹163,251.63
MSE	₹236346207109.23
RMSE	₹486154.51
R^2	0.69

8. Model Deployment using Streamlit

To make our model interactive and usable, we implemented it using **Streamlit**. This allows users to enter car details and receive instant price predictions without needing a complex backend setup.

8.1 Saving the Model using Joblib

We saved the trained model using *Joblib* so that it could be loaded without retraining: This allows us to quickly reuse the model for prediction.

```
import joblib # Library to save/load models

# Save the trained KNN model
joblib.dump(knn, "knn_car_price_model.pkl")

# Save the StandardScaler used for feature scaling
joblib.dump(scaler, "scaler.pkl")

# Save the LabelEncoders for brand & model
joblib.dump(label_encoders, "label_encoders.pkl")

print("Model, scaler, and encoders saved successfully!")
```

8.2 Building the UI with Streamlit

Using *Streamlit*, we developed a simple web app where users can input car details and receive price predictions in real time.

Code for the Streamlit UI:

```
import streamlit as st
import pandas as pd
import numpy as np
import joblib

# Loading the saved model, scaler, and encoders
knn = joblib.load("knn_car_price_model.pkl")
scaler = joblib.load("scaler.pkl")
label_encoders = joblib.load("label_encoders.pkl")
input_data = np.array([[brand_encoded, model_encoded, vehicle_age, km_driven]])
input_data = scaler.transform(input_data)

predicted_log_price = knn.predict(input_data)[0]
predicted_price = np.exp(predicted_log_price)

final_price = apply_depreciation(predicted_price, vehicle_age)

return f"Estimated Price: ₹{final_price:,.2f}"

# Streamlit UI
st.title("Car Price Estimator")
st.subheader("Know Your Car's Worth Before You Sell!")
st.write("Enter the details to get an estimated resale value.")

brand = st.selectbox("Car Brand", label_encoders['brand'].classes_)
model = st.selectbox("Car Model", label_encoders['model'].classes_)
vehicle_age = st.slider("Vehicle Age (Years)", 0, 15, 5)
km_driven = st.number_input("Kilometers Driven", min_value=0, value=50000)

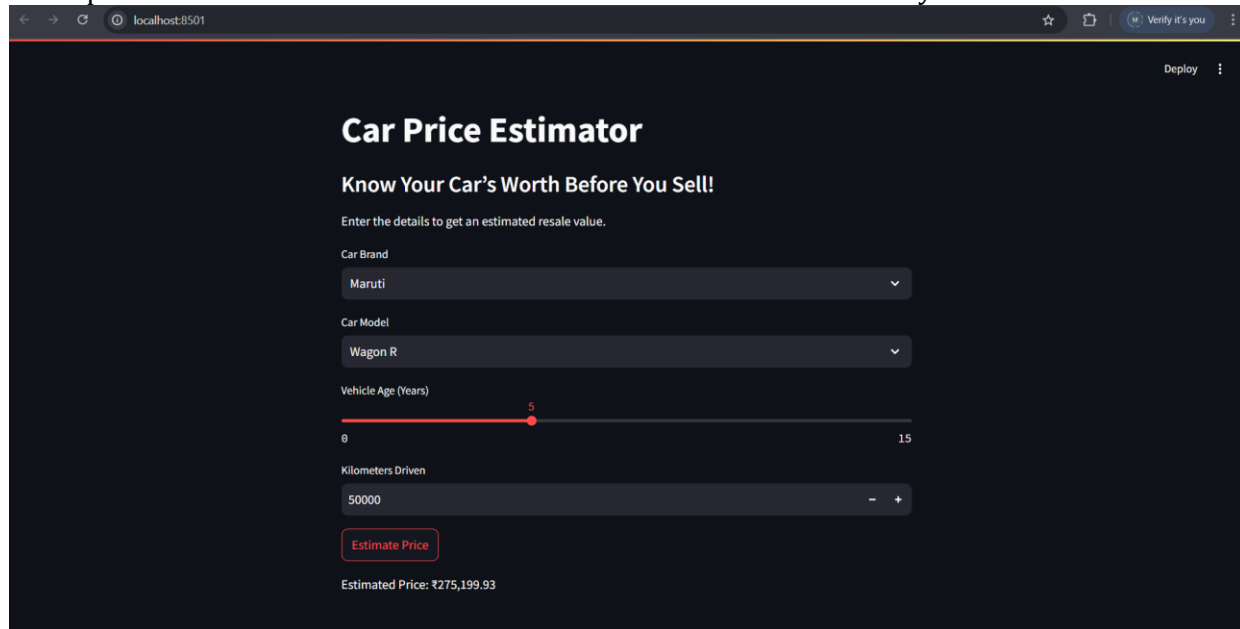
if st.button("Estimate Price"):
    result = predict_car_price(brand, model, vehicle_age, km_driven)
    st.write(result)
```

8.3 Running the Streamlit App

To launch the app locally, we use the following command in terminal:

```
python -m streamlit run "C:\MMD Public\NIRMA\SEM-2\AI-ML\Assignment\predictor.py"
```

This opens a web-based UI where users can interact with the model easily.



9. Future Enhancements and Real-World Integration

To make this model more practical, we can integrate additional real-world factors such as:

- **Location-Based Pricing** – Car prices vary by city, so adding geo-based data would improve accuracy.
- **Market Trends** – Incorporating recent sale trends can make predictions more dynamic.
- **Condition Assessment** – Including vehicle condition (scratches, accidents, servicing history) can refine predictions.
- **User Preferences** – Customizing predictions based on buyer trends and demand fluctuations.

Additionally, this model can be integrated into a real-world system where:

- It connects with **used car marketplace** to provide instant price estimates.
- Dealers can use it to set competitive prices for vehicles.
- Buyers can use it to determine fair market value before making a purchase.

9. Findings and Conclusion

- The model performed well with an R^2 score of 0.89, indicating 89% accuracy.
- Features like year, engine power, and mileage play a major role in determining the selling price.
- More data and feature engineering could improve the accuracy further.

10. References

- Scikit-learn Documentation – <https://scikit-learn.org/>
- CarDekho Dataset – <https://www.kaggle.com/datasets/manishkr1754/cardekho-used-car-data>
- Joblib Documentation - <https://joblib.readthedocs.io/en/stable/>
- Streamlit Tutorials - <https://docs.streamlit.io/develop/tutorials>
- Machine Learning with Python – Research Papers & Online Tutorials